

Group 31 Modeling Notebook

- Ana Margarida Valente, nr 20240936
- Eduardo Mendes, nr 20240850
- Julia Karpienia, nr 20240514
- Marta Boavida, nr 20240519
- Victoria Goon, nr 20240550

Table of Contents

- [0. Import](#)
- [1. Import Datasets](#)
 - [1.1 Encoding Target](#)
- [2. Sampling Datasets](#)
 - [2.1 Undersampling](#)
 - [2.2 Oversampling](#)
- [3. Model](#)
 - [3.1. Logistic Regression](#)
 - [3.2. Decision Tree](#)
 - [3.3. K Nearest Neighbors](#)
 - [3.4. Gaussian Naive Bayes](#)
 - [3.5. Neural Networks](#)
- [4. Ensemble Models](#)
 - [\[4.1. Boosted Models\]](#)
 - [4.1.1. RandomForest](#)
 - [4.1.2. XGBoost](#)
 - [4.1.3. Gradient Boosted Decision Trees](#)
 - [\[4.2. BaggingClassifier Models\]](#)
 - [4.2.1. Bagging DT](#)
 - [4.2.2. Bagging Logistic Regression](#)
 - [4.3. StackingClassifier Models](#)
- [5. Kaggle Submission](#)

0. Import Packages

```
In [1]: ## Import standard data processing libraries  
import numpy as np  
import pandas as pd  
import seaborn as sns  
  
# Importing encoder  
from sklearn.preprocessing import LabelEncoder
```

```

## Import models
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.neural_network import MLPClassifier

## Import ensemble models
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.ensemble import BaggingClassifier
from sklearn.ensemble import StackingClassifier
import xgboost as xgb

# Import Cross Validation methods
from sklearn.metrics import accuracy_score, precision_score, recall_score, f
from sklearn.model_selection import RandomizedSearchCV

# Import imbalanced data methods
from imblearn.over_sampling import SMOTE

# Settings
sns.set()
pd.set_option('display.max_rows', None) # Show all rows
pd.set_option('display.max_colwidth', None) #Show all columns

## Supress warnings
import warnings
warnings.filterwarnings('ignore')

```

```

In [2]: # Code to import imbalanced-learn package for xgboost

# import sys
# !{sys.executable} -m pip install imbalanced-learn

```

1. Import Datasets

Importing the datasets that were created and saved from the preprocessing dataset.

```

In [3]: ## Import the training, validation, and test datasets that were saved from
# train_data = pd.read_csv("train_encoded.csv", low_memory=False)
# validation_data = pd.read_csv("validation_encoded.csv", low_memory=False)
# test_data = pd.read_csv("test_encoded.csv")

# Import the training, validation, and test datasets that were saved from t
train_data = pd.read_csv("train_encoded_std_agg_out.csv", low_memory=False)
validation_data = pd.read_csv("validation_encoded_std_agg_out.csv", low_memc
test_data = pd.read_csv("test_encoded_std_agg_out.csv")

In [4]: # Set the Claim Identifiers as the index
train_data = train_data.set_index("Claim Identifier")

```

```
validation_data = validation_data.set_index("Claim Identifier")
test_data = test_data.set_index("Claim Identifier")
```

```
In [5]: # Seperate target variable from the features in both train and validation
X_train = train_data.drop('Claim Injury Type', axis = 1)
y_train = train_data['Claim Injury Type']

X_val = validation_data.drop('Claim Injury Type', axis = 1)
y_val = validation_data['Claim Injury Type']
```

1.1 Encode Target Variable

Label Encoder for target variable (training and validation):

(This needs to be done in both the preprocessing notebook as well as this modeling notebook to be able to interpret the results properly when a model is tested with the KaggleSubmission csv.)

```
In [6]: #Initiate Label encoder
label_encoder = LabelEncoder()

#Fit the encoder on the training target variable
Y_train_encoded = label_encoder.fit_transform(y_train)

#Transform the training and validation target variable
Y_val_encoded = label_encoder.transform(y_val)

# create a copy of the unencoded target to use when assessing the data - make
y_val_unencoded = y_val.copy()

#Convert the results back to DataFrames while overriding the previous variables
y_train = pd.DataFrame(Y_train_encoded, columns=['encoded_target'], index=y_train.index)
y_val = pd.DataFrame(Y_val_encoded, columns=['encoded_target'], index=y_val.index)
```

2. Sampling Techniques

Below is the WCB dataset class distribution:

1. CANCELLED: 8723
2. NON-COMP: 203607
3. MED ONLY: 48132
4. TEMPORARY: 100810
5. PPD SCH LOSS: 33570
6. PPD NSL: 2759
7. PTD: 49
8. DEATH: 328

Class 7 is significantly lower than the others which will cause the models to not fully train. The following cells is to help address this issue.

Undersampling

The below code takes the size of the minority class + half of the minority class and randomly creates subsamples of each of the majority classes along with all of the minority class and creates a new dataset from the amalgamation of those datapoints.

```
In [7]: ## add the encoded variables back to the x set
training_data_undersampled = pd.concat([X_train, y_train], axis=1)

## Separate majority and minority classes
majority_classes = {}
for x in range(0,8):
    if x != 6:
        majority_classes[x] = training_data_undersampled[training_data_undersampled["encoded_target"] == x]

minority_class = training_data_undersampled[training_data_undersampled["encoded_target"] == 6]

size = int(len(minority_class) + (len(minority_class) * 0.5))

print(size)

## Perform undersampling
undersampled_majority_0 = majority_classes[0].sample(n=size, random_state=0)
undersampled_majority_1 = majority_classes[1].sample(n=size, random_state=1)
undersampled_majority_2 = majority_classes[2].sample(n=size, random_state=2)
undersampled_majority_3 = majority_classes[3].sample(n=size, random_state=3)
undersampled_majority_4 = majority_classes[4].sample(n=size, random_state=4)
undersampled_majority_5 = majority_classes[5].sample(n=size, random_state=5)
undersampled_majority_7 = majority_classes[7].sample(n=size, random_state=7)
undersampled_majority.head()
balanced_data = pd.concat([undersampled_majority_0, undersampled_majority_1, undersampled_majority_2, undersampled_majority_3, undersampled_majority_4, undersampled_majority_5, undersampled_majority_7, minority_class, undersampled_majority_7])

## Separate features and target
X_train = balanced_data.drop(columns='encoded_target')
y_train = balanced_data['encoded_target']

## Check class distribution after undersampling
print("Class distribution after undersampling:", y_train.value_counts())
```

Oversampling

The below cell takes the minority class(es) and created synthetic data for it to match the remaining majority class(es). It uses SMOTE (Synthetic Minority Oversampling Technique).

[https://imbalanced-](https://imbalanced-learn.org/stable/references/generated/imblearn.over_sampling.SMOTE.html)

[learn.org/stable/references/generated/imblearn.over_sampling.SMOTE.html](https://imbalanced-learn.org/stable/references/generated/imblearn.over_sampling.SMOTE.html)

```
In [8]: # Initialize SMOTE
# sampling_strategy=auto : 'not majority' - specifies the class targeted by
# k_neighbors default set to 5. Tried 3 and 10, but 5 had best f1 scores.
smote = SMOTE(sampling_strategy='auto', random_state=42) #k_neighbors=3

# Fit and resample the dataset
X_train, y_train = smote.fit_resample(X_train, y_train)

print("Class distribution after oversampling each at:", y_train.value_counts)
```

Class distribution after oversampling each at: 203607

3. Model

Type of Problem: Multiclassification Problem

Metric used:

As a classification problem, we observed the following metrics to determine the effectiveness of our model: - accuracy - precision - recall - macro f1 score

Each point is measured in a different and observing them all allows us to get an accurate view of our model's results.

```
In [9]: # Functions to help display metrics for all models

# helper method for score_model - not to be used seperately
def print_scores(per_class):
    for x,y in zip(per_class, np.unique(y_val_unencoded)):
        if str(y) == "7. PTD": # add an extra tab for better alignment
            print "["+str(y)+":      \t\t" + str(round(x,2))
        else:
            print "["+str(y)+":      \t" + str(round(x,2))

# displays the scores for Precision, Recall, and F1
def score_model(y_actual, y_predicted, score_train, score_test):

    print("----- F1 -----")
    f1_per_class = f1_score(y_actual, y_predicted, average=None)
    print_scores(f1_per_class)#, y_actual)
    f1_per_weighted = f1_score(y_actual, y_predicted, average='macro')
    print("\nMacro f1: " + str(round(f1_per_weighted, 3)) + "\n")

    print("----- Individual Score Comparisons ----- ")
    print("Train Score: " + str(score_train))
    print("Test Score: " + str(score_test))
    diff = np.abs(score_train - score_test)
    print("Difference: " + str(diff))
```

```

print("----- Accuracy -----\\n")
acc_score = accuracy_score(y_actual, y_predicted)
print("Accuracy Score: " + str(acc_score) + "\\n")

print("----- Precision -----")
precision_per_class = precision_score(y_actual, y_predicted, average=None)
print_scores(precision_per_class)#, y_actual)
precision_weighted = precision_score(y_actual, y_predicted, average='macro')
print("\\nMacro precision: " + str(round(precision_weighted, 3)) + "\\n")

print("----- Recall -----")
recall_per_class = recall_score(y_actual, y_predicted, average=None)
print_scores(recall_per_class)#, y_actual)
recall_per_weighted = recall_score(y_actual, y_predicted, average='macro')
print("\\nMacro recall: " + str(round(recall_per_weighted, 3)) + "\\n")

```

Logistic Regression

Grid Search - Logistic Regression:

```

In [10]: # param_grid = {'C': [0.1, 1, 10], 'solver': ['lbfgs'], 'class_weight': [None]}
# grid_search = GridSearchCV(LogisticRegression(multi_class='multinomial', solver='lbfgs'), param_grid, cv=5)
# grid_search.fit(X_train_std_scaler_encoded, Y_train_encoded_df)

# print("Best Parameters:", grid_search.best_params_)
# print("Best Score:", grid_search.best_score_)

```

Model - Logistic Regression:

```

In [11]: # Create the model
lr_model = LogisticRegression(multi_class='multinomial', solver='lbfgs', C=1)

# Fit the model to the training set
lr_model.fit(X_train, y_train)

# Determine the scores for the model for both train and validation sets
score_train = lr_model.score(X_train, y_train)
score_test = lr_model.score(X_val, y_val)

# Use the model to predict on the validation set
lr_y_pred = lr_model.predict(X_val)

# Display the model metrics using the score_model function
score_model(y_val, lr_y_pred, score_train, score_test)

```

```

----- F1 -----
[1. CANCELLED]:      0.38
[2. NON-COMP]:      0.84
[3. MED ONLY]:      0.28
[4. TEMPORARY]:      0.58
[5. PPD SCH LOSS]:   0.62
[6. PPD NSL]:        0.13
[7. PTD]:            0.01
[8. DEATH]:          0.15

```

Macro f1: 0.373

```

----- Individual Score Comparisons -----
Train Score: 0.7454735102427716
Test Score: 0.6398599368205891
Difference: 0.10561357342218258
----- Accuracy -----

```

Accuracy Score: 0.6398599368205891

```

----- Precision -----
[1. CANCELLED]:      0.25
[2. NON-COMP]:      0.89
[3. MED ONLY]:      0.25
[4. TEMPORARY]:      0.82
[5. PPD SCH LOSS]:   0.54
[6. PPD NSL]:        0.07
[7. PTD]:            0.0
[8. DEATH]:          0.08

```

Macro precision: 0.365

```

----- Recall -----
[1. CANCELLED]:      0.73
[2. NON-COMP]:      0.8
[3. MED ONLY]:      0.31
[4. TEMPORARY]:      0.45
[5. PPD SCH LOSS]:   0.73
[6. PPD NSL]:        0.66
[7. PTD]:            0.17
[8. DEATH]:          0.79

```

Macro recall: 0.58

DECISION TREE

Gridsearch - decision tree:

```

In [12]: ## Initialize the Decision Tree Classifier
         # dt_classifier = DecisionTreeClassifier(random_state=42)

         ## Define the parameter grid to search

```

```
# param_grid = {
#     'criterion': ['gini', 'entropy'],
#     'splitter': ['best', 'random'],
#     'max_depth': [None, 10, 20, 30],
#     'min_samples_split': [2, 5, 10],
#     'min_samples_leaf': [1, 2, 4],
#     'max_features': [None, 'sqrt', 'log2'],
#     'max_leaf_nodes': [None, 10, 20, 30],
#     'min_impurity_decrease': [0.0, 0.1, 0.2]
# }

# # Initialize GridSearchCV:
# grid_search = GridSearchCV(estimator=dt_classifier, param_grid=param_grid,

# # Fit GridSearchCV on the training data
# grid_search.fit(X_train, y_train)

# print("Best Parameters:", grid_search.best_params_)
# print("Best Score:", grid_search.best_score_)

# best_model = grid_search.best_estimator_

# #Best Parameters: {'criterion': 'gini', 'max_depth': 10, 'max_features': 1
# #Best Score: 0.7769977245887005
```

Model - Decision Tree:

```
In [13]: # Test History
# 0.366
# (oversampling) - 0.343 - 21s

# Initialize the Decision Tree Classifier
decision_tree = DecisionTreeClassifier(
    criterion='gini',
    max_depth=10,
    max_features=None,
    max_leaf_nodes=None,
    min_impurity_decrease= 0.0,
    min_samples_leaf= 1,
    min_samples_split=2,
    splitter='best',
    random_state=42
)

# Train the model
decision_tree.fit(X_train, y_train)

# Determine the scores for the model for both train and validation sets
score_train = decision_tree.score(X_train, y_train)
score_test = decision_tree.score(X_val, y_val)

# Make predictions
dt_y_pred = decision_tree.predict(X_val)
```



```
# Display the model metrics using the score_model function
score_model(y_val, dt_y_pred, score_train, score_test)
```

```
----- F1 -----
[1. CANCELLED]:      0.44
[2. NON-COMP]:      0.86
[3. MED ONLY]:      0.22
[4. TEMPORARY]:     0.61
[5. PPD SCH LOSS]:  0.54
[6. PPD NSL]:       0.11
[7. PTD]:          0.01
[8. DEATH]:        0.17
```

Macro f1: 0.369

```
----- Individual Score Comparisons -----
Train Score: 0.7271330307897076
Test Score:  0.6579717550868717
Difference:  0.06916127570283592
----- Accuracy -----
```

Accuracy Score: 0.6579717550868717

```
----- Precision -----
[1. CANCELLED]:      0.33
[2. NON-COMP]:      0.88
[3. MED ONLY]:      0.29
[4. TEMPORARY]:     0.76
[5. PPD SCH LOSS]:  0.44
[6. PPD NSL]:       0.06
[7. PTD]:          0.0
[8. DEATH]:        0.1
```

Macro precision: 0.358

```
----- Recall -----
[1. CANCELLED]:      0.67
[2. NON-COMP]:      0.84
[3. MED ONLY]:      0.17
[4. TEMPORARY]:     0.5
[5. PPD SCH LOSS]:  0.7
[6. PPD NSL]:       0.67
[7. PTD]:          0.1
[8. DEATH]:        0.67
```

Macro recall: 0.541

K Nearest Neighbors

Grid Search - KNN

```
In [14]: ## Define the parameter grid for Randomized Search
         # param_distributions = {
```

```
# 'n_neighbors' : [5,10],
# 'leaf_size': [30, 50],
# 'metric': ['euclidean', 'manhattan'],
# }

# # Initialize RandomizedSearchCV with KNN classifier
# random_search = RandomizedSearchCV(
#     estimator=KNeighborsClassifier(),
#     param_distributions=param_distributions,
#     n_iter=5,
#     cv=2,
#     scoring='f1_macro',
#     verbose=2,
#     random_state=42,
#     n_jobs=-1
# )

# # Fit the random search to the training data
# random_search.fit(X_train, y_train)

# print("Best Parameters:", random_search.best_params_)
# print("Best Score:", random_search.best_score_)

# # Best Parameters: {'n_neighbors': 5, 'metric': 'euclidean', 'leaf_size':
# # Best Score: 0.328445945439427
```

Model - KNN

- KNN will be commented out
- KNN is not appropriate for too large datasets. Too computational expensive due to memorization requirements.
- KNN takes too long to process due to our large dataset.

```
In [15]: # Test History
# # 0.334
# # Initialize the KNN Classifier
# knn_model = KNeighborsClassifier(n_neighbors=5, leaf_size=30, metric='eucl

# # Train the model
# knn_model.fit(X_train, y_train)

# # Determine the scores for the model for both train and validation sets
# score_train = knn_model.score(X_train, y_train)
# score_test = knn_model.score(X_val, y_val)

# # Predict on the validation set
# y_pred = knn_model.predict(X_val)

# # Display the model metrics using the score_model function
# score_model(y_val, y_pred, score_train, score_test)
```

Model - Gaussian Naive Bayes:

- Will be commented out
- Assumes normality, independence, Homogeneity of Variance (Homoskedasticity):

```
In [16]: ## create the model
# model = GaussianNB()

## fit the model to the training set
# model.fit(X_train, y_train)

## determine the scores for the model for both train and validation
# score_train = model.score(X_train, y_train)
# score_test = model.score(X_val, y_val)

## use model to predict on validation set
# y_pred = model.predict(X_val)

## display the model metrics
# score_model(y_val, y_pred, score_train, score_test)
```

Neural Network (MLPClassifier):

GridSearch - MLPClasssifer:

```
In [17]: ## Define the parameter grid
# param_grid = {
#     'hidden_layer_sizes': [
#         (50),          # Larger single-layer model
#         (50, 30),      # Moderate two-layer model
#         (100, 50),     # Larger two-layer model
#         (128, 64, 32)  # Complex three-layer model
#     ],
#     'activation': ['relu', 'logistic'],
#     'solver': ['adam', 'sgd'],
#     'alpha': [0.0001, 0.001],
#     'learning_rate': ['adaptive', 'invscaling']
# }

## Initialize the Neural Network model
# mlp = MLPClassifier(random_state=42)

## Initialize Random Search for hyperparameter tuning
# random_search = RandomizedSearchCV(
#     estimator=mlp,
#     param_distributions=param_grid, # Using param_distributions for random search
#     n_iter=10, # Number of random combinations to try
#     cv=2, # 3-fold cross-validation
#     scoring='f1_macro', # Evaluation metric
#     verbose=2, # Display progress logs
#     n_jobs=-1, # Use all available processors for parallel computation
#     random_state=42 # For reproducibility
```

```
# )

# # Fit the randomized search to the training data
# random_search.fit(X_train, y_train)

#
# print("Best Parameters:", random_search.best_params_)
# print("Best Score:", random_search.best_score_)

# # Best Parameters: {'solver': 'adam', 'learning_rate': 'adaptive', 'hidden
# # Best Score: 0.4192846184298069
```

MODEL - MLPClassifier:

```
In [ ]: # with agreement reached - 0.389 - 14m 57s
# (hidden_layer_sizes=(13,), max_iter=500, random_state=42) - 0.395 (no over
# (hidden_layer_sizes=(15,), max_iter=500, random_state=42) - 0.407 (no over
# (hidden_layer_sizes=(20,), max_iter=500, random_state=42) - 0.407 (no over
# (hidden_layer_sizes=(10,), max_iter=500, random_state=42) - 0.389 (no over
# (hidden_layer_sizes=(10,), max_iter=1000, random_state=42) - 0.389 (no over

# # Initialize the Neural Network model
model = MLPClassifier(hidden_layer_sizes=(50, 30),
                      activation='logistic',
                      solver='adam',
                      alpha=0.0001,
                      learning_rate='adaptive',
                      max_iter=200,
                      random_state=42)

# Train the model
model.fit(X_train, y_train)

# Determine the scores for the model for both train and validation sets
score_train = model.score(X_train, y_train) # Accuracy on training data
score_test = model.score(X_val, y_val)      # Accuracy on validation data

# Use the model to predict on the validation set
y_pred = model.predict(X_val)

# Display the model metrics using the score_model function
score_model(y_val, y_pred, score_train, score_test)
```

```

----- F1 -----
[1. CANCELLED]:      0.51
[2. NON-COMP]:      0.86
[3. MED ONLY]:      0.28
[4. TEMPORARY]:     0.76
[5. PPD SCH LOSS]:  0.65
[6. PPD NSL]:       0.13
[7. PTD]:           0.04
[8. DEATH]:         0.3

```

Macro f1: 0.441

```

----- Individual Score Comparisons -----
Train Score: 0.8609852559096691
Test Score: 0.7300067360401375
Difference: 0.13097851986953157
----- Accuracy -----

```

Accuracy Score: 0.7300067360401375

```

----- Precision -----
[1. CANCELLED]:      0.41
[2. NON-COMP]:      0.88
[3. MED ONLY]:      0.31
[4. TEMPORARY]:     0.77
[5. PPD SCH LOSS]:  0.58
[6. PPD NSL]:       0.09
[7. PTD]:           0.02
[8. DEATH]:         0.24

```

Macro precision: 0.413

```

----- Recall -----
[1. CANCELLED]:      0.67
[2. NON-COMP]:      0.84
[3. MED ONLY]:      0.25
[4. TEMPORARY]:     0.75
[5. PPD SCH LOSS]:  0.75
[6. PPD NSL]:       0.24
[7. PTD]:           0.1
[8. DEATH]:         0.4

```

Macro recall: 0.501

Ensemble Models

Random Forest

Fits a number of decision tree classifiers on various sub-samples of the dataset and uses averaging to improve the predictive accuracy and control over-fitting.

```
In [19]: # 0.379
# (oversampling) - 0.433 (overfitting w/ 0.23 diff) - 6m 5s

# Initialize the Random Forest model
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)

# Train the model
rf_model.fit(X_train, y_train)

# Determine the scores for the model for both train and validation sets
score_train = rf_model.score(X_train, y_train) # Accuracy on training data
score_test = rf_model.score(X_val, y_val)      # Accuracy on validation data

# Use the model to predict on the validation set
rf_y_pred = rf_model.predict(X_val)

# Display the model metrics using the score_model function
score_model(y_val, rf_y_pred, score_train, score_test)
```

```

----- F1 -----
[1. CANCELLED]:      0.6
[2. NON-COMP]:      0.9
[3. MED ONLY]:      0.2
[4. TEMPORARY]:     0.79
[5. PPD SCH LOSS]:  0.64
[6. PPD NSL]:       0.02
[7. PTD]:           0.0
[8. DEATH]:         0.37

```

Macro f1: 0.44

```

----- Individual Score Comparisons -----
Train Score: 0.9999981582165642
Test Score: 0.7855674533122735
Difference: 0.21443070490429073
----- Accuracy -----

```

Accuracy Score: 0.7855674533122735

```

----- Precision -----
[1. CANCELLED]:      0.67
[2. NON-COMP]:      0.87
[3. MED ONLY]:      0.43
[4. TEMPORARY]:     0.72
[5. PPD SCH LOSS]:  0.65
[6. PPD NSL]:       0.14
[7. PTD]:           0.0
[8. DEATH]:         0.65

```

Macro precision: 0.518

```

----- Recall -----
[1. CANCELLED]:      0.54
[2. NON-COMP]:      0.94
[3. MED ONLY]:      0.13
[4. TEMPORARY]:     0.88
[5. PPD SCH LOSS]:  0.63
[6. PPD NSL]:       0.01
[7. PTD]:           0.0
[8. DEATH]:         0.26

```

Macro recall: 0.424

XGBoost

Also using decision trees

```

In [20]: # 0.442
# (oversampling) 0.453 (overfit by 0.10 diff) 4m 18s
# (oversampling) 0.469 (overfit by 0.09 diff) 3m 28s - with agreement reached

# max_depth = 19, n_estimators = 150, lr = 0.6 -> overfitting
xgb_model = xgb.XGBClassifier(
    n_estimators=110, # Number of trees

```

```
learning_rate=0.2, # Step size shrinkage
max_depth=7,      # Maximum depth of a tree
random_state=42,   # For reproducibility
use_label_encoder=False, # Avoid warning for encoding
eval_metric='mlogloss' # Evaluation metric for multi-class classification
)

# Train the model
xgb_model.fit(X_train, y_train)

# Determine the scores for the model for both train and validation sets
score_train = xgb_model.score(X_train, y_train) # Accuracy on training data
score_test = xgb_model.score(X_val, y_val)      # Accuracy on validation data

# Use the model to predict on the validation set
xgb_y_pred = xgb_model.predict(X_val)

# Display the model metrics using the score_model function
score_model(y_val, xgb_y_pred, score_train, score_test)
```



```

----- F1 -----
[1. CANCELLED]:      0.6
[2. NON-COMP]:      0.91
[3. MED ONLY]:      0.18
[4. TEMPORARY]:      0.79
[5. PPD SCH LOSS]:   0.66
[6. PPD NSL]:        0.13
[7. PTD]:            0.0
[8. DEATH]:          0.46

```

Macro f1: 0.467

```

----- Individual Score Comparisons -----
Train Score: 0.8782759188043633
Test Score: 0.786397844467156
Difference: 0.0918780743372073
----- Accuracy -----

```

Accuracy Score: 0.786397844467156

```

----- Precision -----
[1. CANCELLED]:      0.63
[2. NON-COMP]:      0.87
[3. MED ONLY]:      0.46
[4. TEMPORARY]:      0.75
[5. PPD SCH LOSS]:   0.62
[6. PPD NSL]:        0.13
[7. PTD]:            0.0
[8. DEATH]:          0.36

```

Macro precision: 0.477

```

----- Recall -----
[1. CANCELLED]:      0.57
[2. NON-COMP]:      0.95
[3. MED ONLY]:      0.11
[4. TEMPORARY]:      0.84
[5. PPD SCH LOSS]:   0.71
[6. PPD NSL]:        0.13
[7. PTD]:            0.0
[8. DEATH]:          0.66

```

Macro recall: 0.496

Gradient Boosted Decision Trees

```

In [21]: # 16 min = max_depth = 6 - .402 f1

gbdt_model = GradientBoostingClassifier(
    n_estimators=100,      # Number of boosting stages
    learning_rate=0.1,     # Shrinks contribution of each tree
    max_depth=6,           # Limits depth of each tree to prevent overfitti
    random_state=42        # For reproducibility
)

```

```
# Train the model
gbdt_model.fit(X_train, y_train)

# Determine the scores for the model for both train and validation sets
score_train = gbdt_model.score(X_train, y_train) # Accuracy on training data
score_test = gbdt_model.score(X_val, y_val) # Accuracy on validation data

# Use the model to predict on the validation set
gbdt_y_pred = gbdt_model.predict(X_val)

# Display the model metrics using the score_model function
score_model(y_val, gbdt_y_pred, score_train, score_test)
```

```

----- F1 -----
[1. CANCELLED]:      0.59
[2. NON-COMP]:      0.9
[3. MED ONLY]:      0.2
[4. TEMPORARY]:     0.78
[5. PPD SCH LOSS]:  0.66
[6. PPD NSL]:       0.16
[7. PTD]:           0.0
[8. DEATH]:         0.41

```

Macro f1: 0.462

```

----- Individual Score Comparisons -----
Train Score: 0.8648855392987471
Test Score: 0.7791217597324166
Difference: 0.08576377956633052
----- Accuracy -----

```

Accuracy Score: 0.7791217597324166

```

----- Precision -----
[1. CANCELLED]:      0.58
[2. NON-COMP]:      0.87
[3. MED ONLY]:      0.43
[4. TEMPORARY]:     0.75
[5. PPD SCH LOSS]:  0.61
[6. PPD NSL]:       0.13
[7. PTD]:           0.0
[8. DEATH]:         0.3

```

Macro precision: 0.458

```

----- Recall -----
[1. CANCELLED]:      0.59
[2. NON-COMP]:      0.94
[3. MED ONLY]:      0.13
[4. TEMPORARY]:     0.82
[5. PPD SCH LOSS]:  0.71
[6. PPD NSL]:       0.2
[7. PTD]:           0.0
[8. DEATH]:         0.67

```

Macro recall: 0.506

Bagging Decision Tree

```

In [22]: base_model = DecisionTreeClassifier(
            criterion='gini',
            max_depth=10,
            max_features=None,
            max_leaf_nodes=None,
            min_impurity_decrease= 0.0,
            min_samples_leaf= 1,
            min_samples_split=2,

```

```

        splitter='best',
        random_state=42
    )
    bagging_model_mlp = BaggingClassifier(estimator=base_model, n_estimators=100)
    bagging_model_mlp.fit(X_train, y_train)
    bagging_y_pred = bagging_model_mlp.predict(X_val)

    score_train = bagging_model_mlp.score(X_train, y_train)
    score_test = bagging_model_mlp.score(X_val, y_val)

    score_model(y_val, bagging_y_pred, score_train, score_test)

```

```

----- F1 -----
[1. CANCELLED]:      0.45
[2. NON-COMP]:      0.87
[3. MED ONLY]:      0.22
[4. TEMPORARY]:     0.61
[5. PPD SCH LOSS]:  0.54
[6. PPD NSL]:       0.11
[7. PTD]:           0.01
[8. DEATH]:         0.17

```

Macro f1: 0.373

```

----- Individual Score Comparisons -----
Train Score: 0.7313992151546852
Test Score: 0.6640864535910063
Difference: 0.06731276156367894
----- Accuracy -----

```

Accuracy Score: 0.6640864535910063

```

----- Precision -----
[1. CANCELLED]:      0.34
[2. NON-COMP]:      0.88
[3. MED ONLY]:      0.3
[4. TEMPORARY]:     0.77
[5. PPD SCH LOSS]:  0.44
[6. PPD NSL]:       0.06
[7. PTD]:           0.0
[8. DEATH]:         0.1

```

Macro precision: 0.363

```

----- Recall -----
[1. CANCELLED]:      0.67
[2. NON-COMP]:      0.86
[3. MED ONLY]:      0.18
[4. TEMPORARY]:     0.5
[5. PPD SCH LOSS]:  0.71
[6. PPD NSL]:       0.66
[7. PTD]:           0.1
[8. DEATH]:         0.68

```

Macro recall: 0.544

Bagging Logistic Regression

```
In [23]: # 13m - 0.372 - overfit by 0.1 diff
base_model_lr = LogisticRegression(multi_class='multinomial', solver='lbfgs')
bagging_model_lr = BaggingClassifier(estimator=base_model_lr, n_estimators=100)
bagging_model_lr.fit(X_train, y_train)
bagging_y_pred = bagging_model_lr.predict(X_val)

score_train = bagging_model_lr.score(X_train, y_train)
score_test = bagging_model_lr.score(X_val, y_val)

score_model(y_val, bagging_y_pred, score_train, score_test)
```

```

----- F1 -----
[1. CANCELLED]:      0.38
[2. NON-COMP]:      0.84
[3. MED ONLY]:      0.27
[4. TEMPORARY]:      0.59
[5. PPD SCH LOSS]:   0.62
[6. PPD NSL]:        0.13
[7. PTD]:            0.01
[8. DEATH]:          0.15

```

Macro f1: 0.374

```

----- Individual Score Comparisons -----
Train Score: 0.7467326761850035
Test Score: 0.6434834618600762
Difference: 0.10324921432492729
----- Accuracy -----

```

Accuracy Score: 0.6434834618600762

```

----- Precision -----
[1. CANCELLED]:      0.25
[2. NON-COMP]:      0.89
[3. MED ONLY]:      0.25
[4. TEMPORARY]:      0.82
[5. PPD SCH LOSS]:   0.54
[6. PPD NSL]:        0.07
[7. PTD]:            0.0
[8. DEATH]:          0.08

```

Macro precision: 0.365

```

----- Recall -----
[1. CANCELLED]:      0.74
[2. NON-COMP]:      0.8
[3. MED ONLY]:      0.3
[4. TEMPORARY]:      0.46
[5. PPD SCH LOSS]:   0.73
[6. PPD NSL]:        0.67
[7. PTD]:            0.17
[8. DEATH]:          0.79

```

Macro recall: 0.582

Stacking

```

In [24]: # 0.440 - LR -> XGB -> MLP w/ a 0.015 difference in scores (3m 55s)
# 0.425 - LR -> MLP -> XGB w/ a 0.0099 difference in scores (37m)
# 0.410 - MLP -> XGB -> GBC w/ a 0.011 difference in scores (93m 35s)

base_models = [
    ('lr', LogisticRegression(multi_class='multinomial', solver='lbfgs', C=1
    ('dt', DecisionTreeClassifier(
        criterion='gini',

```

```

        max_depth=10,
        max_features=None,
        max_leaf_nodes=None,
        min_impurity_decrease= 0.0,
        min_samples_leaf= 1,
        min_samples_split=2,
        splitter='best',
        random_state=42
    ))
]

nn = MLPClassifier(hidden_layer_sizes=(64, 32), # Two hidden layers: 64 and
                    activation='relu',         # ReLU activation function
                    solver='adam',             # Adam optimizer
                    alpha=0.0001,              # Regularization term (L2)
                    learning_rate_init=0.001,  # Initial learning rate
                    max_iter=200,              # Maximum number of iterations
                    random_state=42)

stacked_model = StackingClassifier(estimators=base_models, final_estimator=nn)
stacked_model.fit(X_train, y_train)
y_pred = stacked_model.predict(X_val)

score_train = stacked_model.score(X_train, y_train)
score_test = stacked_model.score(X_val, y_val)

score_model(y_val, y_pred, score_train, score_test)

```

```

----- F1 -----
[1. CANCELLED]:      0.47
[2. NON-COMP]:      0.85
[3. MED ONLY]:      0.29
[4. TEMPORARY]:      0.67
[5. PPD SCH LOSS]:   0.64
[6. PPD NSL]:        0.13
[7. PTD]:            0.01
[8. DEATH]:          0.31

```

Macro f1: 0.421

----- Individual Score Comparisons -----

Train Score: 0.794949338676961
 Test Score: 0.6775411130725634
 Difference: 0.11740822560439768
 ----- Accuracy -----

Accuracy Score: 0.6775411130725634

```

----- Precision -----
[1. CANCELLED]:      0.35
[2. NON-COMP]:      0.89
[3. MED ONLY]:      0.29
[4. TEMPORARY]:      0.79
[5. PPD SCH LOSS]:   0.56
[6. PPD NSL]:        0.07
[7. PTD]:            0.01
[8. DEATH]:          0.2

```

Macro precision: 0.395

```

----- Recall -----
[1. CANCELLED]:      0.72
[2. NON-COMP]:      0.8
[3. MED ONLY]:      0.28
[4. TEMPORARY]:      0.59
[5. PPD SCH LOSS]:   0.74
[6. PPD NSL]:        0.72
[7. PTD]:            0.07
[8. DEATH]:          0.71

```

Macro recall: 0.579

Weighted Averaging

```

In [25]: ## lr_y_pred_f1 = f1_score(y_val, lr_y_pred, average='macro')
## dt_y_pred_f1 = f1_score(y_val, dt_y_pred, average='macro')
## knn_y_pred_f1 = f1_score(y_val, knn_y_pred, average='macro')
## mplc_y_pred_f1 = f1_score(y_val, mplc_y_pred, average='macro')
## rf_y_pred_f1 = f1_score(y_val, rf_y_pred, average='macro')
## xgb_y_pred_f1 = f1_score(y_val, xgb_y_pred, average='macro')
## gbdt_y_pred_f1 = f1_score(y_val, gbdt_y_pred, average='macro')

## f1_score(y_actual, y_predicted, average='macro')

```



```

# # Assign weights based on F1 scores
# #weights = [lr_y_pred_f1, dt_y_pred_f1, knn_y_pred_f1, mlpc_y_pred_f1, rf_
# weights = [mlpc_y_pred_f1, xgb_y_pred_f1, gbdt_y_pred_f1]
# weights = np.array(weights) / np.sum(weights) # Normalize weights

# # Make weighted predictions
# # lr_probs = lr_model.predict_proba(X_val)[: , 1]
# # dt_probs = decision_tree.predict_proba(X_val)[: , 1]
# # knn_probs = knn_model.predict_proba(X_val)[: , 1]
# # mlpc_probs = mlpc_model.predict_proba(X_val)[: , 1]
# # rf_probs = rf_model.predict_proba(X_val)[: , 1]
# # xgb_probs = xgb_model.predict_proba(X_val)[: , 1]
# # gbdt_probs = gbdt_model.predict_proba(X_val)[: , 1]

# # Aggregate predictions using weights
# weighted_probs = (
#     # weights[0] * lr_probs +
#     # weights[1] * dt_probs +
#     # weights[2] * knn_probs +
#     weights[0] * mlpc_probs +
#     # weights[4] * rf_probs +
#     weights[1] * xgb_probs +
#     weights[2] * gbdt_probs)

# # Final predictions (threshold = 0.5)
# final_predictions = (weighted_probs >= 0.2).astype(int)

# # Evaluate the ensemble
# final_f1 = f1_score(y_val, final_predictions, average='macro')
# print(f"Weighted Ensemble F1 Score: {final_f1:.2f}")

```

5. Kaggle Submission

```

In [26]: # get the model prediction
# y_pred_test = bagging_model_lr.predict(test_data)

In [27]: ## decode the prediction labels back to their original values
# decoded_labels = label_encoder.inverse_transform(y_pred_test)

In [28]: # # # combine the prediction values with their claim identifiers into a dataframe
# kaggle_submission = pd.DataFrame({"Claim Identifier": test_data.index, "Claim Status": decoded_labels})
# kaggle_submission.head()

In [29]: # # Compile the resulting dataframe into a csv file named "Kaggle_submission.csv"
# # this will be found in the directory the file is currently running from
# # if a file exists with the same name, it will overwrite it with the new dataframe
# kaggle_submission.to_csv("Kaggle_Submission.csv", index=False)

```